

likelihoods

GK

February 4, 2019

Now we will estimate weights by using eBayes approach. To do so, we will optimize the log-likelihood function using `optim()` function. More precisely, we will find values of weights which maximize the log-likelihood function.

As a **first step**, we have to define the log-likelihood function:

$$L(w) = f(Y|w) = \prod_{i=1}^n f(y_i|w)$$

since we assume that y_i are independent for $i = 1, \dots, n$ Hence,

$$L(w) = \left(\frac{1}{(2\pi)^n \det(\Sigma)}\right)^{1/2} \exp\left(-\frac{1}{2}(y'\Omega^{-1}y)\right)$$

where Ω is function of weight: w_i .

Given the likelihood function above, we know that maximizing the log of likelihood is the same as maximizing the likelihood. So, we will stick with log-likelihood function from there.

As a **next step**, we will add penalty term to the log-likelihood function to smooth it and make estimation more efficient.

We add a penalty on a function as structure of weight parameter defined in CC-Process 1c. Hence from our goal is to maximize log of the following joint likelihood function:

$$p(w, y) = p(y|w)p(w)$$

where $p(y|w)$ is a data model, i.e $L(w)$ and $p(w)$ is already defined in CC-Process 1c s.t

$$p(y|w) = N(y; 0, \Omega(X, w))$$

$$p(w) = N(w; 0, \Pi)$$

Note that this idea comes from LASSO or Ridge regression methods of parameter estimation.

We can write the function in R as follows:

```
loglkl_mvn_penalty <- function(w,d) { #

  time <- ncol(d)-2
  # calculate the covariance matrix for yw
  omega <- gcalc_corr(d,w)

  # calculate the covariance matrix for w
  sigma <- 1/(0.36)*calc_Sigma(time+1, 0.99)

  # data model: log_likelihood
  p_yw <- dmvnorm(d[,ncol(d)], mean = rep(0, nrow(d)),
                 sigma = omega, log = TRUE)
```

```

# prior on w: log_likelihood
p_w <- dmvnorm(w, mean = rep(0, time+1),
               sigma = sigma , log = TRUE)

# the posterior which will be maximized
result <- p_yw + p_w

return(result)
}

```

In pair with the following function:

```

estimate_w <- function(opt_f, grad, pars, d, maxit){
# Estimates w using eBayes approach, i.e finds MLE
# estimates of w.
#
# Args:
#   opt_f: likelihood function to be optimized
#   grad:  gradient of opt_f
#   pars:  initial values for w
#   d:     observed data
#   maxit: maximum number of iterations
#
# Output:
#   opt:   results of optim()

  opt <- optim(par = pars, opt_f, d = d,
               control = list(fnscale = -1,maxit=maxit),
               gr = grad)

  # print true values of w
  print("True values of w")
  print(w)

  # estimated w
  return(opt)
}

```

Note that we have added gradient of log-likelihood function f for better optimization. Since the differenttiation of the original likelihood function is way difficult, we tried first numerical (approximate) version of gradient:

$$\nabla f = (f'_{w_0}, \dots, f'_{w_T})$$

where $w = (w_0, \dots, w_T)$. By the fundamentals of calculus we have that:

$$f'_{w_i}(w_0, \dots, w_T) = \lim_{h \rightarrow \infty} \frac{f(w_0, \dots, w_i + h, \dots, w_T) - f(w_0, \dots, w_T)}{h}$$

We have implemented this in practice as follows:

```
calc_gradient_num <- function(f,w,d,epsilon=10^-8){
# Calculates the gradient of a function numerically
#
# Args:
#   f: function (log-likelihood function)
#   d: data frame, last column is response Y,
#       others are input X's
#   w: weights, in a vector form
#
# Output:
#   gr: gradients, in a vector form

  n <- length(w)
  gr <- numeric(n)
  for(i in 1:n) {
    h <- rep(0,n); h[i] <- epsilon
    gr[i] <- (f(w+h,d)-
              f(w,d))/epsilon
  }
  return(gr)
}
```

Here h is epsilon, given as an infinitesimally small number.

Taking into account assumption about density of weights, log-likelihood function and its gradient, we obtained estimates of w which are reasonably close to the true weight parameters. However, there is a small concern about the estimates with different initial values used in `optim()`, so `optim()` might be converging to the local maxima. To see that, we can plot 3D graphs of log-likelihood using the functions given in *3D_plot_loglkl.R* (!!! More work: attach images, additionally code has to be refined)

To better optimize the likelihood function, we will modify some assumptions about the model (specifically, weight parameter). We will transform w_t into the log-scale s.t

$$\log(w) \sim N(-1 \text{ or } -2, \Sigma)$$

Following the simple rules of variable transformation, we have the following first result:

1. For $w = (w_0, \dots, w_T)$ define $\theta = (\theta_0, \dots, \theta_T)$ s.t $\theta_i = \log(w_i)$. Hence, $\theta = \log(w) = h(w)$ s.t $h^{-1}(\theta) = e^\theta = w$. Then

$$p_\theta(\theta) = p_w(h^{-1}(\theta)) | \det J_{h^{-1}(\theta)} |$$

$$J_{h^{-1}(\theta)} = \text{diag}(e^{\theta_0}, \dots, e^{\theta_T})$$

Hence,

$$|\det J_{h^{-1}(\theta)}| = \prod_{i=0}^T e^{\theta_i} = e^{\sum_{i=0}^T \theta_i} = e^{(T+1)\bar{\theta}}$$

$$p_{\theta}(\theta) = (2\pi)^{-(T+1)/2} (\det \Pi)^{-1/2} e^{(T+1)\bar{\theta}} e^{-\frac{1}{2}(e^{\theta_0}, \dots, e^{\theta_T}) \Pi^{-1} (e^{\theta_0}, \dots, e^{\theta_T})'}$$

Note that covariance matrix Π used in the density formula is already defined in CC-Process 1c (for reference). We will sample θ from the distribution given above and convert sampled θ 's to w : the variable of interest. (For this purpose we would like to employ MCMC method (using package 'mcmc' in R) to sample from not known density. However, it requires the output of density to be **scalar**, but we have a vector)

Sampling multivariate r.v from unknown distribution can be done in two (for me) ways: 1. By employing Metropolis-Hasting algorithm/Block-Wise sampling method (here I had a difficulty in finding the appropriate proposal distribution of the same dimension as a target distribution) 2. By employing "rstan" in R (I will probably try this method later, once I get acquainted with Stan)

Sample code is given in *supplementary_functions.R*. Here is sample code for approach 1:

```
## this function does not work right now
simulate_log_w <- function(t){
# Simulates log(weights) from transformed MVN ~ (GP)
# given in CC-Process 1c
#
# Args:
#   t: Time(T) or length(X)-1
#
# Output:
#   w: weights, in a vector form

# define the density
sigma <- 1/(0.36)*calc_Sigma(t+1, 0.99)
dens <- function(theta){
  # theta is a vector of length T
  result <- (2*pi)^(-(t+1)/2)*(det(sigma))^(1/2)
  *exp(sum(theta))*exp(-1/2*theta%*%solve(sigma)%*%t(theta))
  return(result) ### need a scalar, not vector
}

require(mcmc)
out <- metrop(dens, sigma, 1e+3)
return(out$batch)
```

```

}

## use this website to modify the above code
# We would like to use MH at https://theclevermachine.wordpress.com/tag/multivariate-sa

## Error: <text>:17:9: unexpected '*'
## 16:      result <- (2*pi)^(-(t+1)/2)*(det(sigma))^(1/2)
## 17:      *
##      ^

```

Sample code for approach 2:

```
#STANNNNN
```

2. For $w = (w_0, \dots, w_T)$ define $\theta = (\theta_0, \dots, \theta_T)$ s.t $\theta_i = \log(w_i)$. Here we will stick with the traditional (not proven) assumptions; since we want a density which takes into account high correlation among variables, we will leave covariance matrix structure Π the same as AR(1). Hence,

$$p_{\theta}(\theta) = MVN(\theta; -2, \Pi)$$

Note that mean = -2 can be adjusted (we may choose any in range (-5,-1) probably). We will sample θ from the distribution given above and convert sampled θ 's to w : the variable of interest.

Sample code is given below:

To sample weights wisely we could find a distribution which takes into account facts about the structure of w :

1. Highly correlated weights, with possibly AR(1) structure
2. Concentrated around zero
3. Skewed to the right
4. Positive s.t $0 \leq w \leq 1$

There are some processes that likely to follow these characteristics. One is **Beta AR(1)** process model, which has the following structure: (cited from internet/article: An Autoregressive Process for Beta Random Variables, by E.Mackenzie and Bayesian Model Selection for Beta Autoregressive Processes by R.Casarin)

Also we can improve the optimization process by deriving the real gradient of log-likelihood function with penalty term. For that purpose we will utilize the method given in Rasmussen's book GP for ML. We will take the following steps:

1. Function of interest for which we have to find gradient is log-likelihood with penalty term:

$$\begin{aligned} \log(p(w, y)) &= \log(p(y|w)) + \log(p(w)) \\ &= \log\left(\left(\frac{1}{(2\pi)^n \det(\Omega)}\right)^{1/2} e^{y' \Omega^{-1} y}\right) + \log\left(\left(\frac{1}{(2\pi)^{T+1} \det(\Pi)}\right)^{1/2} e^{w' \Pi^{-1} w}\right) \end{aligned}$$

where Π is constant for a fixed Time(T)

2. Parameter of interest, or the variable by which we have to calculate the gradient is $w = (w_0, \dots, w_T)$ of length $T+1$.
3. We can notice that it is not difficult to derive the gradient, i.e first derivatives of $\log(p(w))$ w.r.t w , since its structure reduced to quadratic polynomial after taking the log.
4. We are interested particularly in

$$\frac{\partial \log(p(w))}{\partial w_k} = 0 + 0 - \frac{1}{2} \frac{\partial (w' \Pi^{-1} w)}{\partial w_k}$$

We take a closer look:

$$\begin{aligned} &\frac{\partial (w' \Pi^{-1} w)}{\partial w_k} = \\ &= -\frac{1}{2} (w_0 \Pi_{00}^{-1} + w_1 \Pi_{10}^{-1} + \dots + w_T \Pi_{T0}^{-1}, \dots, w_0 \Pi_{0T}^{-1} + \dots + w_T \Pi_{TT}^{-1}) (w_0, \dots, w_T)' \\ &\quad - \frac{1}{2} (w_0 w_0 \Pi_{00}^{-1} + w_0 w_1 \Pi_{10}^{-1} + \dots + w_0 w_T \Pi_{T0}^{-1} + w_1 w_0 \Pi_{01}^{-1} + w_1 w_1 \Pi_{11}^{-1} + \\ &\quad \dots + w_1 w_T \Pi_{T1}^{-1} + w_T w_0 \Pi_{0T}^{-1} + \dots + w_T w_T \Pi_{TT}^{-1}) \end{aligned}$$

Taking into account the special structure of the covariance matrix Π , which is AR(1), we know that its inverse - precision matrix will have the following structure:

$$\Pi^{-1} = \begin{bmatrix} \Pi_{00}^{-1} & \Pi_{01}^{-1} & 0 & 0 & \cdot & \cdot & \cdot & 0 \\ \Pi_{10}^{-1} & \Pi_{11}^{-1} & \Pi_{12}^{-1} & 0 & \cdot & \cdot & \cdot & 0 \\ 0 & \Pi_{21}^{-1} & \Pi_{22}^{-1} & \Pi_{23}^{-1} & \cdot & \cdot & \cdot & 0 \\ 0 & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ 0 & 0 & 0 & \cdot & \cdot & \cdot & \Pi_{T-1T-1}^{-1} & \Pi_{TT}^{-1} \end{bmatrix}$$

Here everything is zero except the diagonal and its two parallels right next to it. This will make the calculation of gradient more easier and faster.

Hence, for $k \in \{0, \dots, T\}$

$$\begin{aligned} \frac{\partial \log(p(w))}{\partial w_k} &= -\frac{1}{2} \left(w_{k-1} (\Pi_{k-1,k}^{-1} + \Pi_{k,k-1}^{-1}) \right. \\ &\quad \left. + 2w_k \Pi_{k,k}^{-1} + w_{k+1} (\Pi_{k,k+1}^{-1} + \Pi_{k+1,k}^{-1}) \right) \end{aligned}$$

As an ultimate measure of accuracy we can calculate the MSE for each method of estimation and different initial values. Appropriate lines of codes are given in *assess_accuracy.R* file and their implementation is given in *test_on_simulation.R* file

Last important part of optimization is a time required for the whole run of `optim()` function. Lines of codes for time assessment (in benchmark style) is given in the *assess_runtime.R* file